Duke Compute Cluster Workshop

3/28/2018

Tom Milledge
rc.duke.edu
rescomputing@duke.edu

Outline of talk

- Overview of Research Computing resources
- Duke Compute Cluster overview
- Running interactive and batch jobs
- Using Slurm job arrays
- Running multi-core and parallel jobs
- Specifying job dependencies
- Live demo and questions

Duke Research Computing services & resources

- Cluster computing - Duke Compute Cluster
- Virtual machines - RAPID (Research Toolkits), “Clockworks”
- Data storage - Duke Data Commons (NIH),
- Research computing networks - high speed, Software-Defined Network, “Protected Network”
- Consultation
- Access to national resources (XSEDE, OSG)
- Education & training
- Technology development

The Duke Compute Cluster (DCC)

- Formerly known as the “DSCR”
- 667 compute nodes with approximately 14209 CPU cores (as of March 2018)
- Most nodes are purchased by labs and depts
- Some are provided by the University
- 500 TB of primary storage (Isilon X-series)
- Red Hat Enterprise Linux 7
- Uses the Slurm job queueing system

Accessing the Duke Compute Cluster

- From on-campus (University or DUHS networks):

ssh NetID@dcc-slogin.oit.duke.edu

This will connect to one of the three DCC login nodes

- From off campus, first , use the Duke VPN:

https://oit.duke.edu/net-security/network/remote/vpn/

Copying files and directories

- Use scp or rsync for Linux or Mac workstations
- Use winscp for Windows: https://winscp.net
- Copying a file to the DCC (“push”)
  - `% scp data001.txt netid@dcc-slogin-02.oit.duke.edu:.`
- Copying a file from the DCC (“pull”):

```
% scp netid@dcc-slogin-02.oit.duke.edu:output.txt .
```

- Use either scp -r (small files) or rsync –av (large files)
- Pushing a directory:

```
rsync –av dir1/  netid@dcc-slogin-02.oit.duke.edu:. or
scp -r dir1/ netid@dcc-slogin-02.oit.duke.edu:.
```

- Pulling a directory:

```
rsync –av netid@dcc-slogin-02.oit.duke.edu:~/dir1 .
scp -r netid@dcc-slogin-02.oit.duke.edu:~/dir1 .
```

Duke Compute Cluster file systems

/dscrhome (a symlink to /hpchome/group)

- Primary storage on the Isilon X-series filers
- 250 GB group quota (typical)
- Two week tape backup (TSM)

/work

- Temporary storage for large data files
- Place for temporary files associated with running jobs (I/O)
- Not backed-up
- Subject to purges based on file age and/or utilization levels
- 200 TB total volume size, unpartitioned

/datacommons

- Archival storage on the Isilon NL-series filers
- available for $82/TB/year

Slurm resources

- The DCC web pages

https://rc.duke.edu/the-duke-compute-cluster/

"Official" SLURM docs

http://schedmd.com/slurmdocs

- Older SLURM documentation

https://computing.llnl.gov/linux/slurm/slurm.html

- Comes up a lot in Google searches
- outdated – use schedmd.com instead

Slurm jobs run in “partitions”

- Most DCC partitions are dept-owned machines
- These can only be used by members of the group
- Submitting to a group partition gives “high-priority”
- Submit to partitions with “--partition=” or “-p“, e.g.

  `#SBATCH –p (partition name)` (in a script) or

  `srun –p (partition name)--pty bash –i` (interactively)
- The default DCC partition is called “common”
- The common partition gives “low-priority” to most ESX hosts

DCC Partitions

There are different DCC partitions to which batch jobs and interactive sessions can be directed:

- **common**, for jobs that will run on the DCC core nodes (up to 64 GB RAM).
- **common-large**, for jobs that will run on the DCC core nodes (64-240 GB GB RAM).
- **gpu-common**, for jobs that will run on DCC GPU nodes.
- **Group partitions** (partition names varies), for jobs that will run on lab-owned nodes

Running an interactive job

- Reserve a compute node by typing

```
srun --pty bash -i
```

```
tm103@dscr-slogin-02 ~ $ srun --pty bash -i
srun: job 186535 queued and waiting for resources
srun: job 186535 has been allocated resources
tm103@dscr-core-11 ~ $
```

```
tm103@dscr-encode-11 ~ $ squeue -u tm103
JOBID PARTITION NAME USER ST TIME NODES NODELIST(REASON)
186535 common bash tm103 R 0:14 1 dscr-core-11
```

- I now have an interactive session in the common partition on node dscr-encode-11

SLURM commands

- sbatch

Submit a batch job (like “qsub”)

- #SBATCH

Specify job parameters (like “#$”)

- squeue (like “qstat”)

Show lists of jobs

- scancel (like “qdel”)

Delete one or more batch jobs

- sinfo (like “qhost”)

Show info about machines

- scontrol

Show cluster configuration information

sbatch

- Use “sbatch” (all lower case) to submit

  `sbatch test.q`

- Use “#SBATCH” (upper case) in your scripts for scheduler directives, e.g.

```
#SBATCH --mem=1G
#SBATCH --output=matlab.out
```

- All SLURM directives can be given on the command line instead of the script.
- http://slurm.schedmd.com/sbatch.html

sbatch example

```
#!/bin/bash
#
#SBATCH --output=test.out
uname –n # print hostname
```

This prints the name of the compute node in the file "test.out”

```
tm103@dscr-slogin-02 ~/slurm $ sbatch simple.q
Submitted batch job 186554
tm103@dscr-slogin-02 ~/slurm $ cat test.out
dscr-compeb-14
```

Long-form commands example

```
#!/bin/bash
#SBATCH --output=slurm.out
#SBATCH --error=slurm.err
#SBATCH --mem=100 # 100 MB RAM
#SBATCH --partition=abyss # e.g.
uname -n 1>&2 #prints hostname to the error file
```

- For a user in the “abyss” group, this job will run in high priority on an “abyss” lab node.

Short-form commands example

- SLURM short commands don't use "=" signs

```
#!/bin/bash
#SBATCH -o slurm.out
#SBATCH -e slurm.err
#SBATCH --mem=4G # 4 GBs RAM
#SBATCH -p somelab #only for members
 uname -n 1>&2 #prints hostname to the error file
```

Matlab example script

```
#!/bin/bash
#SBATCH -J matlab
#SBATCH -o slurm.out
#SBATCH --mem=4G # 4 GB RAM
/opt/apps/matlabR2016a/bin/matlab -nojvm -nodisplay -singleCompThread -r my_matlab_program
```

- The "`-singleCompThread`" command is required to prevent uncontrolled multithreading

Slurm memory directives

- **This is a hard limit** – always request a little more

--mem=<*MB*>

- The amount of memory required per node

--mem-per-cpu=<*MB*>

- The amount of memory per CPU core
- For multi-threaded jobs
- Note: --mem and --mem-per-cpu are mutually exclusive

Slurm parallel directives

- All parallel directives have defaults of 1

-N <*number*> How many nodes (machines)

-n <*number*> or **--ntasks**=<*number*> How many parallel jobs (“tasks”)

-c, **--cpus-per-task**=<*ncpus*>

- Use **-c** for multi-threaded jobs
- The **--ntasks** default is one CPU core per task, but the **--cpus-per-task** option will change this default.

Multi-threaded (multi-core) example

```
!/bin/bash
#SBATCH -J test
#SBATCH -o test.out
#SBATCH -c 4
#SBATCH --mem-per-cpu=500 #(500 MB)
myApplication -n $SLURM_CPUS_PER_TASK
```

- The value of $SLURM_CPUS_PER_TASK is the number after “-c”
- This example starts a single, multi-threaded job that uses 4 CPU cores and 2 GB (4x500MB) of RAM

OpenMP multicore example

```
!/bin/bash
#SBATCH -J openmp-test
#SBATCH -o slurm.out
#SBATCH -c 4
export OMP_NUM_THREADS=$SLURM_CPUS_PER_TASK
myOpenMPapp  # will run on 4 CPU cores
```

- This sets **$OMP_NUM_THREADS** to the value of **$SLURM_CPUS_PER_TASK**

Slurm job arrays

- a mechanism for submitting and managing collections of similar jobs
- Job arrays are only supported for batch jobs and the array index values are specified using the **`--array`** or **`-a`** option
- http://slurm.schedmd.com/job_array.html

Matlab job array example

```
#!/bin/bash
#SBATCH --output=matlab.out
#SBATCH --array=1-100
#SBATCH --mem=4G
/opt/apps/MATLAB/R2012b/bin/matlab
-nojvm -nodisplay –singleCompThread
-r "rank=$SLURM_ARRAY_TASK_ID;my_prog;quit”
```

- Start 100 Matlab programs, each with a different “rank”, e.g. 1,2, ... 100

Python job array example

```
#!/bin/bash
#SBATCH -e slurm.err
#SBATCH --array=1-5000
python myCode.py
```

```
$ cat test.py
import os
rank=int(os.environ['SLURM_ARRAY_TASK_ID'])
...
```

- Start 5000 Python jobs, each with a different “rank”, initialized from `SLURM_ARRAY_TASK_ID`

Running MPI jobs

- Supported MPI versions
- Intel MPI
- OpenMPI
- SLURM MPI jobs use “--ntasks=(num)“

https://wiki.duke.edu/display/SCSC/Running+MPI+Jobs

Compiling with OpenMPI

```
tm103@dscr-slogin-02  ~ $ export PATH=/opt/apps/slurm/openmpi-2.0.0/bin:$PATH
tm103@dscr-slogin-02  ~ $ which mpicc
/opt/apps/slurm/openmpi-2.0.0/bin/mpicc
tm103@dscr-slogin-02  ~ $ mpicc -o openhello /dscrhome/tm103/misc/slurm/openmpi/hello.c
tm103@dscr-slogin-02  ~ $ ls -l openhello
-rwxr-xr-x. 1 tm103 scsc 9184 Sep  1 16:08 openhello
```

OpenMPI job script

```
#!/bin/bash
#SBATCH -o openhello.out
#SBATCH -e openhello.err
#SBATCH -n 20
export PATH=/opt/apps/slurm/openmpi-2.0.0/bin:$PATH
mpirun -n $SLURM_NTASKS openhello
```

OpenMPI example output

```
tm103@dscr-slogin-02 ~/misc/slurm/openmpi $ cat
openhello.out
dscr-core-01, rank 0 out of 20 processors
dscr-core-01, rank 1 out of 20 processors
dscr-core-01, rank 2 out of 20 processors
dscr-core-01, rank 3 out of 20 processors
dscr-core-01, rank 4 out of 20 processors
dscr-core-03, rank 13 out of 20 processors
dscr-core-03, rank 14 out of 20 processors
dscr-core-03, rank 10 out of 20 processors
dscr-core-03, rank 11 out of 20 processors
dscr-core-03, rank 12 out of 20 processors
dscr-core-02, rank 8 out of 20 processors
dscr-core-02, rank 9 out of 20 processors
dscr-core-02, rank 5 out of 20 processors
...
```

Intel MPI example

```
#!/bin/bash
#SBATCH --ntasks=20
#SBATCH --output=intelhello.out
export I_MPI_PMI_LIBRARY=/opt/slurm/lib64/libpmi.so
source /opt/apps/intel/intelvars.sh
srun -n $SLURM_NTASKS intelhello
```

GPU nodes

To run a GPU batch job, add the job script lines

#SBATCH -p gpu-common --gres=gpu:1
#SBATCH -c 6

To get an interactive GPU node session, type the command line

srun -p gpu-common --gres=gpu:1 -c 6 --pty bash –i

```
tm103@dscr-slogin-02  ~ $ srun -p gpu-common --gres=gpu:1 -c 6 --pty bash -i
tm103@dscr-gpu-01  ~ $ /usr/local/cuda-7.5/samples/1_Utilities/deviceQuery/
deviceQuery
...
Detected 1 CUDA Capable device(s)

Device 0: "Tesla K80"
  CUDA Driver Version / Runtime Version          7.5 / 7.5
  CUDA Capability Major/Minor version number:    3.7
  Total amount of global memory:                 11520 MBytes (12079136768 bytes)
  (13) Multiprocessors, (192) CUDA Cores/MP:     2496 CUDA Cores
...
```

Job dependencies

- https://hcc-docs.unl.edu/display/HCCDOC/Job+Dependencies
- Start job “dep2” after job “dep1”

```
$ sbatch dep1.q
Submitted batch job 666898
```

- Make a note of the assigned job ID of dep1

```
$ sbatch --dependency=afterok:666898 dep2.q
```

- Job dep2 will not start until dep1 finishes

Job dependencies with arrays

- Wait for specific job array elements

```
sbatch --depend=after:123_4 my.job
sbatch --depend=afterok:123_4:123_8 my.job2
```

- Wait for entire job array to complete

```
sbatch --depend=afterany:123 my.job
```

- Wait for entire job array to complete successfully

```
sbatch --depend=afterok:123 my.job
```

- Wait for entire job array to complete and at least one task fails

```
sbatch --depend=afternotok:123 my.job
```

Live demo notes

- `df –h`
- `srun --pty bash -i`
- `squeue | more`
- `squeue -S S`
- `squeue–S S|grep –v PD`
- `squeue –u (NetID)`
- `sbatch (job script)`
- `scancel (job id)`
- `uname –n, sleep, top`
- `sinfo | grep –v common`
- `scontrol show node (node name)`
- `scontrol show job (job id)`